

# CPSC457 Fall 2020 - Assignment 1

Due date is posted on D2L.

**Individual assignment. Group work is NOT allowed.**

Weight: 15% of the final grade.

## Motivation for this assignment:

- Well written Python code can run faster than badly written C/C++ code.
- Common reason for bad performance is using unnecessary system calls.
- You will attempt to improve the performance of a badly written C++ program.

Please clone the GIT repository: <https://gitlab.com/cpsc457/public/a1> . It contains:

<code>palindrome.py</code>	Python 3 program that reads in text from standard input and reports the longest palindrome to standard output.
<code>slow-pali.cpp</code>	Not a very good C++ version of <code>palindrome.py</code> . Feel free to re-use any part of this code in your solution.
<code>Makefile</code>	Makes compilation a bit easier.
<code>t1.txt ... t5.txt</code>	Some test files.

For this assignment we will use the following definitions:

Standard input	Please read <a href="http://www.linfo.org/standard_input.html">http://www.linfo.org/standard_input.html</a> .
White space	Whatever <code>isspace()</code> reports as white-space. Read the man page for <code>isspace()</code> for more information.
Word	Non-zero-length sequence of characters delimited by white space.
Palindrome	Any word that remains the same after reversing it and ignoring the case. Examples of palindromes: “ <code>Did</code> ”, “ <code>01-!-10</code> ”, “ <code>x</code> ”
Longest palindrome	If there are multiple palindromes, the longest one is reported. If multiple palindromes have the same maximum length, report the first one.

## Q1 - Written question (4 marks)

For this question you will compare the performance of the python program to the C++ program by using the `time` Unix utility. To time how long it takes to execute the python program on different test files, for example `t4.txt` and `t5.txt`, execute the following commands:

```
$ time python3 palindrome.py < t4.txt
Longest palindrome: redder
real    0m0.300s
user    0m0.296s
sys     0m0.003s
$ time python3 palindrome.py < t5.txt
Longest palindrome: Detartrated
real    0m0.023s
user    0m0.016s
sys     0m0.007s
```

Before you can time the C++ program, you need to compile it first. Here is how to do it by hand:

```
$ g++ -O2 -Wall slow-pali.cpp -o slow-pali
```

Here is how to compile the C++ code by using the included Makefile:

```
$ make
g++ -O2 -Wall slow-pali.cpp -o slow-pali
```

Now you can time the resulting executable `slow-pali`:

```
$ time ./slow-pali < t4.txt
Longest palindrome: redder
real    0m2.929s
user    0m1.477s
sys     0m1.450s
$ time ./slow-pali < t5.txt
Longest palindrome: Detartrated
real    0m0.003s
user    0m0.000s
sys     0m0.003s
```

Answer the following questions:

- Time the python code and C++ code on `t4.txt` and `t3.txt` using the `time` utility. Copy/paste the outputs from the terminal output to your report.
- How much time did the C++ and python programs spend in kernel vs user mode?
- Why is the python program faster on some inputs when compared to C++ code? Why is the python program slower on other inputs?

## Q2 - Programming question (15 marks)

It is embarrassing to have C++ perform slower than Python. Your job is to improve the code in `slow-pali.cpp` by writing a new implementation called `fast-pali.cpp`. Your new implementation should be faster than both `palindrome.py` and `slow-pali.cpp` for all possible inputs! If you wish to write C code, write `fast-pali.c` instead.

### Requirements

- Your program must read input from standard input.
- You are only allowed to use the `read()` system call wrapper. You cannot use any other APIs, such as `mmap()`, `fopen()`, `fread()`, `fgetc()`, or C++'s streams. Hint: adjust to code to use `read()` with a buffer size of 1MiB.
- Do not store the entire input in memory. You need to write your code so that it can handle any input size, even if it is bigger than the available memory.
- Your program must run on `linux.cpsc.ucalgary.ca`. Use SSH to test your code!

### Valid input

Your program must be able to handle any text input of up to 2GiB in size. You may assume that no word will be longer than 1024 bytes.

Some test files are available to you, but it is expected that you create your own test files to help you validate your solutions. Your TAs will grade your code on inputs that are not published to you. Only valid inputs will be used to grade your code.

## Q3 - Written question (4 marks)

Measure the performance of your solution from Q2 and compare it to the timings you obtained for `slow-pali.cpp` and `palindrome.py` in Q1. Answer the following:

- a) Is your program from Q2 faster than `slow-pali.cpp`? Why do you think that is?
- b) Is your program faster or slower than `palindrome.py` and why?

Justify your answers by using the `strace` utility with the `'-c'` command line option. Include the output of `strace` in your report.

## Submission

Submit two files for this assignment to D2L:

1. Answers to the written questions combined into a single file called `report.pdf`. Do not use any other file format.
2. Your solution to Q2 in a file called `fast-pali.cpp` or `fast-pali.c`.

## General information about all assignments:

1. All assignments are due on the date listed on D2L. Late submissions will be not be marked.
2. Extensions may be granted only by the course instructor.

3. After you submit your work to D2L, verify your submission by re-downloading it.
4. You can submit many times before the due date. D2L will simply overwrite previous submissions with newer ones. It is better to submit incomplete work for a chance of getting partial marks, than not to submit anything. Please bear in mind that you cannot re-submit a single file if you have already submitted other files. Your new submission would delete the previous files you submitted. So please keep a copy of all files you intend to submit and resubmit all of them every time.
5. Assignments will be marked by your TAs. If you have questions about assignment marking, contact your TA first. If you still have questions after you have talked to your TA, then you can contact your instructor.
6. All programs you submit must run on [linux.cpsc.ucalgary.ca](http://linux.cpsc.ucalgary.ca). If your TA is unable to run your code on the Linux machines, you will receive 0 marks for the relevant question.
7. Unless specified otherwise, you must submit code that can finish on any valid input under 10s on linux.cpsc.ucalgary.ca, when compiled with `-O2` optimization. Any code that runs longer than this may receive a deduction, and code that runs for too long (about 30s) will receive 0 marks.
8. **Assignments must reflect individual work.** For further information on plagiarism, cheating and other academic misconduct, check the information at this link: <http://www.ucalgary.ca/pubs/calendar/current/k-5.html>.
9. Here are some examples of what you are not allowed to do for individual assignments: you are not allowed to copy code or written answers (in part, or in whole) from anyone else; you are not allowed to collaborate with anyone; you are not allowed to share your solutions with anyone else; you are not allowed to sell or purchase a solution. This list is not exclusive.
10. We will use automated similarity detection software to check for plagiarism. Your submission will be compared to other students (current and previous), as well as to any known online sources. Any cases of detected plagiarism or any other academic misconduct will be investigated and reported.